

Research Article

#Earthquake: Twitter as a Distributed Sensor System

Andrew Crooks

*Department of Computational
Social Science
George Mason University*

Arie Croitoru

*Center for Geospatial Intelligence
and Department of Geography and
Geoinformation Science
George Mason University*

Anthony Stefanidis

*Center for Geospatial Intelligence
and Department of Geography and
Geoinformation Science
George Mason University*

Jacek Radzikowski

*Center for Geospatial Intelligence
and Department of Geography and
Geoinformation Science
George Mason University*

Abstract

Social media feeds are rapidly emerging as a novel avenue for the contribution and dissemination of information that is often geographic. Their content often includes references to events occurring at, or affecting specific locations. Within this article we analyze the spatial and temporal characteristics of the *twitter* feed activity responding to a 5.8 magnitude earthquake which occurred on the East Coast of the United States (US) on August 23, 2011. We argue that these feeds represent a hybrid form of a sensor system that allows for the identification and localization of the impact area of the event. By contrasting this with comparable content collected through the dedicated crowdsourcing ‘Did You Feel It?’ (DYFI) website of the U.S. Geological Survey we assess the potential of the use of harvested social media content for event monitoring. The experiments support the notion that people act as sensors to give us comparable results in a timely manner, and can complement other sources of data to enhance our situational awareness and improve our understanding and response to such events.

Address for correspondence: Andrew Crooks, Department of Computational Social Science, George Mason University. Research I, CSC Suite, Level 3, 4400 University Drive, Fairfax, VA 22030, USA. E-mail: acrooks2@gmu.edu

1 Introduction

Enabled by Web 2.0, ubiquitous computing, and corresponding technological advancements, social media has drastically altered the concepts of information contribution, dissemination, and exchange (Kaplan and Haenlein 2010). By empowering the general public to publish and distribute user-generated content, social media has enabled them to bypass the need to utilize centralized, authoritative agencies (e.g. news organizations). In a sense, this makes virtually every citizen a potential contributor or user of information. The power of social media to communicate information of societal importance has been demonstrated very effectively during the last few years. One could argue that a defining moment in this direction was the coordinated use of Twitter, Facebook, and YouTube by Iranian protesters in June 2009 to communicate to a global audience the state of unrest in the streets of Tehran, bypassing the ruling regime's restrictions on information dissemination (Newsweek 2009). A comparable situation on a larger geographic scale was witnessed in early 2011 during the Arab Spring events in North Africa (Christensen 2011, Howard et al. 2011). More recently social media has been heavily involved during the Occupy Wall Street protests in the U.S. (Huffington Post 2011, Christian Science Monitor 2011). In a sense, such technology enables the crowds to publicly broadcast information, on a scale never seen before.

The rise in Web 2.0 technology has also enabled crowdsourcing (Howe 2008), the delegation of tasks traditionally performed by professionals to the community at large. We can identify numerous geographic applications that benefit from this. For example, car navigation HD TrafficTM from TomTom (www.tomtom.com/hdtraffic) provides real time traffic information, updated every three minutes, while the Waze (<http://www.waze.com/>) application allows a community of users to exchange geographically explicit information such as traffic incidents in real time. In both instances the quality of information depends on the level of participation and the number of users. In the related area of public transportation, citizens are crowdsourcing information on bus arrival time through Tiramisu (Steinfeld et al. 2011). Furthermore, crowds are used to detect potholes in roads (see for example, SeeClickFix (<http://www.seeclickfix.com>) or FixMyStreet (<http://www.fixmystreet.com>) or to report a variety of local problems ranging from poor road conditions and graffiti to background noise levels in their neighborhoods, by using GPS-enabled smart phones (Demirbas et al. 2010). A sensor-driven variant of these applications is Street Bump (<http://www.appbrain.com/app/street-bump-%28alpha%29/com.citizapps.streetbump>) which takes advantage of accelerometers and GPS sensors embedded in smart phones to automatically capture and report street pothole and bump information. Conceptually, this falls under the 'Participatory Sensing' framework, whereby ad-hoc sensor networks are formed taking advantage of sensors embedded in mobile devices, to enable public and professional users to gather, analyze and share local information (Burke et al. 2006). Currently there exists a variety of applications whereby sensor measurements are crowdsourced, including Japan Geigermap (<http://japan.failedrobot.com>) that visualizes crowdsourced radiation Geiger counter readings from across Japan, or the contribution of weather data in Weather Underground (<http://www.wunderground.com>). Participatory sensing still adheres to the Volunteered Geographic information (VGI, Goodchild 2007) principle of the general public contributing on purpose particular pieces of geographic information, most frequently towards defined tasks.

Like the crowdsourcing systems presented above, social media feeds often convey *geographic* information, as people frequently comment on events happening at or affecting their location, or refer to locations that represent momentary social hotspots (e.g. by referring to the location of a protest, or to the area hit by a natural disaster). Nevertheless, this represents a deviation from the well-established concept of crowdsourcing in geographic applications (Fritz et al. 2009, Goodchild and Glennon 2010): unlike Wikimapia or OpenStreetMap, social media feeds are not a vehicle for citizens to explicitly and purposefully contribute geographic information to update or expand geographic databases. Instead, the geographic content is embedded in the contributors' comments, and has to be harvested and analyzed before it can be used. This type of geographic information that can be harvested from social media feeds could be referred to as Ambient Geographic Information (AGI, Stefanidis et al. 2012) and represents an extension of the concept of VGI. One could argue that whereas VGI is primarily crowdsourcing, with specific tasks outsourced to the public at large, AGI is crowdharvesting, with the general public broadcasting information that can be harvested in a meaningful manner.

The proliferation of social media feeds and corresponding AGI content fosters the emergence of micro-blogging as a new type of a distributed sensor system, with citizen bloggers acting as sensors, and their comments (e.g. in the form of tweets) conveying relevant information, often with a corresponding geographic footprint. In this article we offer an analysis of the performance of micro-blogging as a sensor system for geographic event detection by considering the Mineral, VA earthquake of August 23, 2011 as a sample case. Some early publications in this emerging research direction have addressed the use of Twitter content to generate map mashups to support collaborative real-time mapping (Field and O'Brien 2010), and provided an overview of emerging opportunities for harvesting geospatial content (Stefanidis et al. 2012). Furthermore, Earle et al. (2010) attempted assessing how fast tweeters reacted to the smaller (4.3 magnitude) and much more localized earthquake of Morgan Hill, CA in March, 2009. Our work in this article will extend this emerging research direction, by analyzing the manner in which earthquake information was reported in the blogosphere, in order to identify spatial and temporal characteristics of this information dissemination avenue. This is a critical step towards assessing the value of harvesting geospatial information from social media feeds for event detection and analysis. We consider in particular Twitter as a representative sample social media environment for our work. In our analysis we will also compare Twitter reaction to the event (which represents a form of AGI) to the VGI content contributed by the public through the dedicated 'Did You Feel It?' (DYFI) website of the U.S. Geological Survey (USGS). DYFI builds upon the work of Wald et al. (1999) and represents a prototypical crowdsourcing website (<http://earthquake.usgs.gov/earthquakes/dyfi/>), with a focus on earthquakes. It provides a web portal to collect citizen responses to earthquakes, and integrates this information to model earthquake activity (Atkinson and Wald 2007). By investigating a large earthquake (5.8 magnitude) with a substantial impact area in the East Coast of America, by analyzing the spatial and temporal characteristics of the Twitter community response to it, and by comparing this to crowdsourced information, our work extends our understanding of the value of this novel type of geospatial information.

The article is organized as follows. In Section 2 we make a case for an analogy between social media feeds and sensor systems. In Section 3 we discuss information harvesting and analysis from social media feeds, followed by a discussion on the

geolocation content of tweets in Section 4. In Section 5 we provide an assessment of the role of micro-blogging as a geosensor system, using the Virginia earthquake data. Finally, in Section 6 we offer our outlook assessment.

2 Social Media Feeds as Sensor Systems

Quite interestingly, early work to link Twitter and sensor systems started from the opposite direction than the one we pursue. For example, Demirbas et al. (2010) presented a system architecture to publish sensor observations over Twitter. In the same direction, the UK snow-map Twitter mashup (<http://uksnowmap.com/>) crowdsources snow coverage in the UK, with Twitter contributors using a particular tweet format (containing a specific hashtag and a location in the form of a postcode) in order to report snow fall in their vicinity. That information is then aggregated to generate snow coverage maps. Again, this is an example of crowdsourcing that simply makes use of Twitter to contribute particular information.

However, the argument that we present here is its reverse, namely that Twitter is a form of a hybrid sensor system. When citizens communicate information through micro-blogging they act in many ways just like a sensor would: something is catching their attention, and they report it. Unlike sensors though, this information typically is not provided directly in the form of a quantified measurement (e.g. ‘42.5 degrees Fahrenheit’), but it is often hidden within the text of a message (e.g. ‘yesterday was a cold day in London’). Furthermore, unlike typical sensors that always operate on specific bands of the spectrums, or collect specific types of measurements, humans operate in a wide range of the socio-cultural spectrum, commenting in one message on a natural phenomenon, and in the next on a political issue. Yet despite these slight variations, human bloggers operate in a manner comparable to sensors in a sensor network.

If we consider the typical architecture of individual sensor nodes within wireless sensor networks (WSN), we can see that a prototypical node comprises certain key elements (Akyildiz et al. 2002) as shown in Figure 1:

- One or more dedicated sensors to collect local data for monitoring the specific types of phenomena that the WSN is designed to observe;
- A microcontroller processing unit to control and execute the basic functionalities of the sensor;




Figure 1 Components of a prototypical sensor node

- A transceiver to communicate with other nodes in the WSN and thus ensure information flow within the network;
- Data storage capability, for future retrieval and analysis; and of course,
- An energy source to support its long term operation.

In a micro-blogging analogy to the physical sensor node, the human blogger is the sensor and the microcontroller: she collects the information that she considers remarkable and important to communicate (Sternberg 2009), in order to document one's life, provide a commentary, or to build a community (Nardi et al. 2004, Java et al. 2009, Naaman et al. 2010), and publishes this information. In this example, the transceiver is the actual micro-blogging application (e.g. Twitter) that enables the dissemination of this information among the other nodes of the network (the remainder of the Twitter community). The presence of an energy source is not important in this particular situation, as users use a variety of modalities, ranging from handhelds to laptops and desktops, which are presumably free of the limitations presented when nodes are deployed in remote, often inaccessible locations (Selavo et al. 2007). Lastly, data storage is handled by the micro-blogging site itself, as this is where the data are actually stored and reside.

In comparison to standard WSN deployments, where the number of nodes and the area they cover are both rather limited, micro-blogging offers massive numbers of nodes, and practically worldwide distribution. In the spring of 2011, Twitter announced that it had over 200 million accounts, distributed all over the world. Among these accounts, it is estimated that Twitter has 100 million active users, logging in at least once a month, and 50 million users who do so daily. Furthermore, its #numbers entry (<http://blog.twitter.com/2011/03/numbers.html>) indicates that a record 572,000 new accounts were created on a single day (March 12, 2011, the day after the Fukushima earthquake and resulting nuclear disaster), while an average of 140 million tweets are sent daily, resulting in a billion tweets sent every week. As a measure of reference it is worth mentioning that Facebook has stated that 600 million people visit it each month, with half of them using it daily.

However, a key distinction between micro-blogging and wireless sensor networks is the lack in the former of an explicit application-serving cooperative framework that is a characteristic of the latter (Akyildiz et al. 2002). While there is interaction between Twitter users (e.g. through retweeting or following) and this can lead to the establishment of geographically defined complex social network groups (see for example, Stefanidis et al. 2012), the collaboration of these groups tends to be ad-hoc and not as rapid as in traditional wireless sensor networks (Honey and Herring 2009, Zhao and Rosson 2009). Nevertheless, as we assert later in this article, micro-blogging emerges as a hybrid form of a distributed sensor system, supporting key operations like the effective monitoring of natural events.

3 Harvesting and Analyzing Social Media Feeds

Harvesting information from social media feeds is in essence a web-mining process (Kosala and Blockeel 2000, Sakaki et al. 2010, Russell 2011). It entails in general three operations: extracting data from the data providers (various social media servers) via application programming interfaces (APIs); parsing, integrating, and storing these data in a resident database; and then analyzing these data to extract information of interest.

There exist a number of tools that perform parts of these processes, such as 140 kit (<http://140kit.com/>) or Twapperkeeper, (<http://twapperkeeper.com/>) but these are limited in their scalability with respect to large datasets. Sites such as Ushahidi (<http://www.ushahidi.com/>) also provide a means to collect social media feeds and disseminate this information over the web. However, currently available tools offer limited capabilities to add context to content, or to support detailed analysis, thus forcing the development of custom systems to perform the abovementioned three operations.

Original social media feeds can be retrieved from source data providers through queries. This entails submitting a query in the form of an http request and receiving in response data in XML format (e.g. Atom or RSS). The query parameters may be for example, based on location (e.g. specifying an area of interest to which the feed is related), time (e.g. specifying a period of interest), content (e.g. specifying keywords or hashtags), or even user handle/ID. In response to these queries, and depending on the characteristics of the information provided by the service, we can receive from the server just metadata or metadata and actual data. A representative example of the first case is Flickr, where the query result contains exclusively metadata information (e.g. author, time, and geolocation when available), and information on how to access the actual image itself. Twitter is a representative example of the second, where the data received in response to a query are actual tweets and associated metadata (e.g. user information, time of tweet publication, geolocation when available, and information on whether this particular tweet is in response to or retweet of an earlier message).

Once this information is harvested from the social media server it can be parsed to become part of a local database (e.g. implemented using PostgreSQL), thus creating a local mirror of the content of the original server for the entries specified by our queries. Depending on the subject, the queries may be periodic, or may be intensified during episodes of crisis. While the information harvested from social media in this manner is not explicitly geographic, it does include implicit geospatial content, thus rendering it suitable for geographic analysis.

The diverse content of these datasets and of user-contributed web-content in general has introduced novel applications for event detection and contribution validation. In the work that is most closely related to our analysis, Hyvarinen and Saltikoff (2010) focused on imagery contributed to Flickr, and addressed the quality of the provided metadata in terms of its descriptive content (confirming for example that over 70% of the contributed imagery labeled as hail actually depicts hailstorm), positional accuracy (assessing it to be close to 1–2 km for landmark imagery), and temporal accuracy (less than 15 minutes for 90% of all contributed imagery, and substantially less for automatically populated metadata, e.g. when contributions are made using handheld devices). They also performed a very limited analysis of Flickr imagery using two images of hailstorms and confirming them with radar data for the area and time they depict.

The temporal, spatial, and social dynamics of Twitter activity during a major forest fire were addressed by De Longueville et al. (2009) in an effort to investigate the performance of Location-Based Social Networks (LBSN) as information creation and dissemination platforms. In the same direction, an analysis of Twitter activity from people on the ground during two emergency events (grass fires and river flooding) indicated the possibility of increased situational awareness through it (Vieweg et al. 2010). This issue was also addressed by Singh et al. (2010) who introduced the concept of social pixels and social imagery as visualizations of patterns of user interest. Furthermore,

MacEachren et al. (2011) presented a geovisual analytics approach to aggregate social Twitter-derived information in place-time-concept indexing schemes, for enhanced situational awareness.

Social media feeds have been used quite extensively for event detection from Web and Social Media (WSM) information (Brownstein et al. 2009, Sakaki et al. 2010, Okasaki and Matsuo 2012). Notable applications addressed the early detection of contagious outbreaks by monitoring influenza-related blogging trends during the emergence of the U.S. 2008 flu season (Corley et al. 2010), and the analysis of the structure of the social network and corresponding relationships to improve our prediction of the spread of these outbreaks (Christakis and Fowler 2010). In the same application domain of bioinformatics, Ginsberg et al. (2009) used web-contributed information, in the form of search engine query data, to detect influenza epidemic outbreaks, and showed that their approach beat the standard methods used by the Centers for Disease Control to monitor these outbreaks.

4 Geolocation Content of Twitter Feeds

Geolocation information in tweets can be provided directly by the contributing bloggers, if they decide to make this information available, or it can be deduced from IP addresses using any of the IP geolocation solutions (see Eriksson et al. 2010, Poese et al. 2011). In this article we are focusing on geolocation information that is contributed either directly by the user or provided through the client application.

This geolocation information may be available either in the form of precise coordinates as shown in Figure 2, or in a descriptive manner (e.g. listing a city name, as shown in Figure 3). It is typically harvested from Twitter using the capabilities provided by the communication protocol linking the client to it. For example, the World Wide Web Consortium (W3C) Geolocation API enables scripting code to access device information from web browsers of any web-capable device (e.g. a mobile phone or a laptop; see Doty and Wilde 2010). In this way information is collected in a dynamic mode, reflecting the actual location from where a tweet was sent. In addition to this, geolocation information can also be harvested from the content of users' profiles, but this is less reliable as it is static and does not necessarily reflect user location at the moment that the tweet was sent.

Statistics on the percentage of geolocatable tweets vary from one part of the world to another, depending on the prevalence of GPS-equipped smartphones, and to cultural acceptance of the location information option in Twitter (Java et al. 2009, Miller and Wortham 2010, Hecht et al. 2011). Early on Java et al. (2009) had reported that approximately half of Twitter users in their study provided some location information in the corresponding entry of their profiles (either coordinates or description). More recently Cheng et al. (2010) reported that 5% of users in their study listed actual coordinates, with another 21% of users listing locational information at the city level, while Hecht et al. (2011) reported that two out of three users in their study provided such geolocation information. Stefanidis et al. (2012) reported that approximately 16% of the Twitter feeds in their experiments had detailed location information with it in the form of coordinates, while another 45% of the tweets they collected had some geolocation information at coarser granularity (e.g. the city level). These variations in the availability of geolocation information tend to be geographic in nature, as providing precise

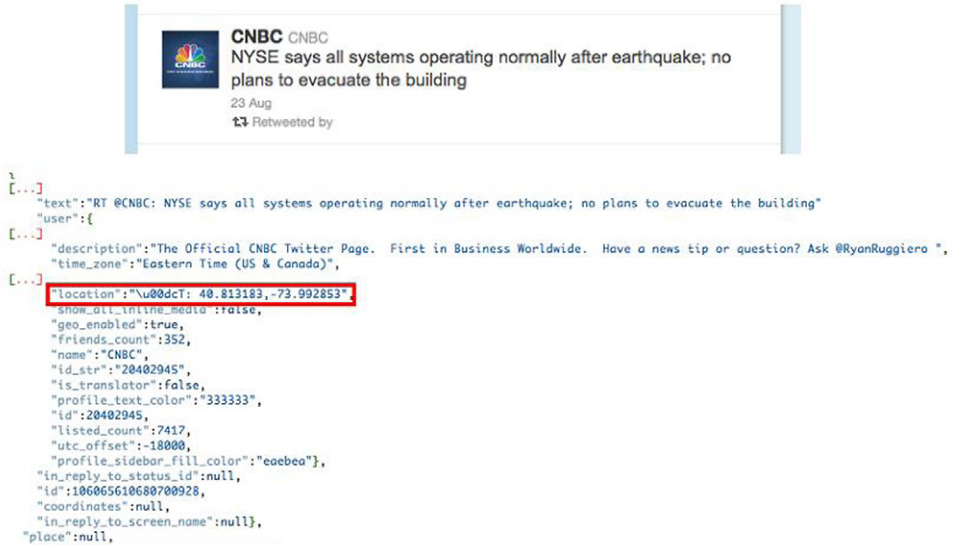


Figure 2 Example of a geolocated tweet. At the *top* we see the tweet as it appears on a follower's stream, and this is the Twitter message that is actually displayed in the browser. At the *bottom* we see highlights from the information retrieved through the search API for this particular tweet, including its geolocation information (marked by the box) in the form of precise coordinates. The coordinates correspond to a location in the Fairview borough of New Jersey

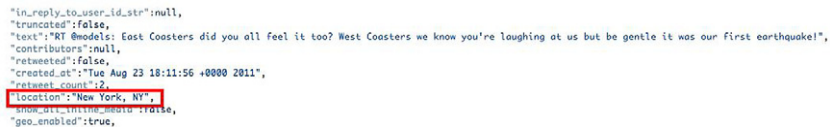


Figure 3 An example of descriptive geolocation information recovered for another tweet similar to the one presented in Figure 2

geolocational information is common practice when using mobile devices to post. Accordingly, one can reasonably argue that precise geolocation information will be more frequently available in areas where the latest mobile technology is more easily and rapidly adopted.

Attempting to assess the granularity of descriptive geolocation information, studies in the social media community indicate that the large majority of users who provide such information tend to gravitate towards listing their city as their location (Ludford et al. 2007, Hecht et al. 2011). Furthermore, recent work by Cheng et al. (2010) showed that even without any IP or geolocation information provided, users' location could often be estimated at the city level based purely on the content of their messages. Comparable work has been performed on predicting the location of Flickr users through an analysis of their contributions' content (Popescu and Grefenstete 2010), and of Facebook users through an analysis of their social network (Backstrom et al. 2010).

5 Twitter Response to the Mineral, VA Earthquake

In order to assess the performance of micro-blogging and AGI as a sensor system for geographic event detection we use a major earthquake as a test case. The August 23, 2011 earthquake that is the subject of this analysis had an epicenter at 37.936° N, 77.933° W, in the rural Louisa County (VA), 8 km SSW of Mineral, VA (135 km SW of Washington, DC). This magnitude 5.8 earthquake struck at 17:51:04 UTC time (13:51 local time), and was the largest earthquake to hit the Eastern part of the U.S. since 1944. According to the USGS this was the most widely felt earthquake in U.S. history (http://www.dmmme.virginia.gov/DMR3/va_5.8_earthquake.shtml) as it impacted the very densely populated East Coast of the U.S., and was felt from central Georgia in the South to central Maine in the North, and from the East Coast all the way to Illinois in the West.

The Mineral, VA earthquake is a particularly suitable event for our analysis as it is the largest magnitude earthquake to hit the eastern half of the U.S. in more than half a century. Accordingly, it has a more substantial impact on the locals than comparable events in earthquake-prone areas like California, making the local society more reactionary to this event. Furthermore, the event affected a large geographic area, including major metropolitan centers, thus providing a substantial sample to perform our analysis.

5.1 Dataset and Data Streaming Rates

We analyze the response to this earthquake in Twitter by harvesting a 1% random sample of Twitter feeds directly from the Twitter API¹ in the period immediately following this event, and we mine the data to find references to the earthquake through keyword (earthquake or earth and quake) and hashtag search (#earthquake or #quake). Hashtags represent a community-adopted means to add context to tweets by using them to label terms of particular significance. Practically, they serve as user-driven tools through which real-life events are assigned a brief blogosphere reference term.

We are interested in an assessment of the performance of Twitter feeds as a novel type of geosensor network (Nittel et al. 2008) and accordingly we consider only geolocated tweets. Furthermore, we are interested in tweets initiated within the contiguous 48 states and thus filtered out any contributions from outside this region's latitude and longitude extent. From our 1% sample of tweets over a period of 8 hours immediately following the earthquake we have collected 144,892 tweets with a reference to earthquake, out of which 21,362 were geolocated tweets with precise coordinate information (for an average of 2,670 such geolocated tweets per hour) and this is the data corpus we use in this analysis. In this particular dataset the precisely geolocated tweets represent 14.7% of the tweets with a reference to the earthquake. It is worth mentioning that of these geolocated tweets, 1,558 (7.3%) were retweets, a process that is typically faster than composing a new one.

Table 1 provides a summary of the manner in which this information was streaming in Twitter. It is worth noting that the first tweet in our data corpus arrived 54 seconds after the event (at 17:51:58 UTC) from location Laurel, VA (approximately 50 Km from the epicenter) and its content was: "Earthquake in Richmond". In comparison, Earle et al. (2010) reported that the first tweet after the March 30, 2009, Morgan Hill, CA earthquake arrived only 19 seconds after that event. An argument can

Table 1 A list of geolocated Twitter reports of earthquake during the first hour of the earthquake (from among our data corpus). The first column lists time period in minutes and seconds. The second column lists the number of precisely geolocated tweets referring to the event at distinct time intervals that are part of our data corpus (and the rate per minute of these tweets). The third column lists the cumulative number of such tweets (from the event time)

Time period (min., sec.)	Geolocated earthquake tweets		Cumulative number of tweets
	Number	Rate per min.	
00m01s – 01m00s	2	2	2
01m01s – 02m00s	91	91	93
02m01s – 03m00s	221	221	314
03m01s – 04m00s	281	281	595
04m01s – 05m00s	324	324	919
05m01s – 10m00s	1142	228	2061
10m01s – 60m00s	11,668	233	13,729

be made that the faster response recorded in Morgan Hill is explained by the fact that Californian residents are very familiar with earthquakes and thus use their past experience of comparable events to respond quicker to new incidents, whereas the East Coast public is not as familiar and its response is slower. Work in psychology fully supports this notion, as it has been shown that repeated exposure to a stimulus leads to the development of processing fluency for this particular stimulus (Nessler et al. 2005), making our brain identify it faster, a process that eventually leads to habituation (Yamaguchi et al. 2004). Such enhanced memory mechanisms have also been shown to be activated in response to the frequent exposure to otherwise traumatic or critical incidents (Dyregrov et al. 2000) as would be the case with the earthquake tremors.

Within 2 minutes of the event we approached 100 tweets, while within 5 minutes we are approaching 1,000. As our data corpus is from a random 1% sample of Twitter it is reasonable to expect that these numbers should be multiplied by 100 if we are to estimate the overall response to the event in the Twitter community, raising these figures to close to 10,000 tweets within 2 minutes, and close to 100,000 within 5 minutes which, as we will see in Section 6, is comparable to the number of reports contributed to the USGS DYFI website over a period of nearly 8 hours after the earthquake.

5.2 Reaction Time and Response Patterns as Function of Distance from Epicenter and Impact Area Detection

While Table 1 shows the temporal trend in geolocated data streaming from Twitter, it does not convey information on the originating locations of these tweets. In this section we investigate temporal patterns of Twitter response as a function of distance from the event epicenter. Figure 4 (top) shows patterns of reaction time of tweet responses over the first 400 seconds (6 minutes 20 seconds) after the earthquake. The horizontal axis shows distance from the epicenter (expressed as km), and the vertical shows the Twitter




Figure 4 Response pattern as function of distance from epicenter for the first 400 seconds after the earthquake. At the *top* we see a plot of *reaction time*, *versus distance* of all tweets during that period. At the *bottom* we show the histogram of the number of tweets as a function of angular distance

community reaction time (in seconds). Reaction time is defined as the time lapsed between the actual event and a tweet referring to it, and is an indicator of the community reaction to this particular event. Each dot in this plot represents a particular tweet in our data corpus, with a corresponding (time, distance) set of coordinates. At the top of that graph we have listed major cities corresponding to a particular distance, to provide a more intuitive reference frame.

A visual inspection of Figure 4 shows that we can easily identify a dense cluster of tweets on the left-hand side of the graph, extending from the epicenter all the way to approximately 700 to 950 km away from it (the area between Boston and Chicago in the graph). We have marked the furthestmost point of this threshold region by a red line in the figure for easy reference. The same observation can be made in the histogram of tweets as shown in the bottom of Figure 4, which records the number of tweets as a function of distance from the epicenter. In that histogram too we observe a substantial drop as we move beyond this point, which we consider to be a threshold in Twitter data stream activity. The spatiotemporal patterns of tweet data streams before and after that threshold are clearly differentiable. Before it (to the left of the red line) we have the overwhelming majority of relevant tweets, resulting in very high density in space and time, whereas beyond it (to the right of it) we have sporadic references only. In the area before the threshold we have large groups of responses from the major metropolitan areas (Washington, D.C., Philadelphia, New York City), which manifest themselves as dark identifiable strands in the graph, indicating nearly continuous flow of information from these locations. Similarly, at the corresponding parts of the histogram (Figure 4, bottom) we identify peaks. At the same time, we have a sufficient number of responses in between these metropolitan areas, to provide good spatial coverage over that area. Beyond the threshold, there is no discernible pattern of data streams as the feeds are separated from each other in space and time.

In order to decipher the meaning of this pattern of data streams we took the locations of the 40 tweets that fall in the threshold sliver between approximately 700 and 950 km of distance from the epicenter in Figure 4 (roughly the dots presented are between the distance spot marked Boston and the one marked Chicago in this figure) and plotted them on top of the USGS Community Decimal Intensity (CDI) map as shown in Figure 5. The CDI map shows the estimated perceived intensity of the earthquake at different locations, and also visualizes the extent of its impact area. As we see the tweets (shown as circles in Figure 5) from this threshold sliver fall at the edge of what is identified to be the earthquake's impact area. Accordingly, we observe that by analyzing the patterns of data streaming (and without any additional information) we can derive an excellent approximation of the earthquake's impact area within as few as 400 seconds of the event, which can be very critical for disaster response.

The two clearly differentiable patterns of activity inside and outside the impact area can be referred to as response to the physical event versus response to the cyber event (its presence in and coverage by the media). One could argue that the first is the reaction to the earthquake, whereas the latter is the response to #earthquake, the corresponding cyber reference term to the news story. From a sensor network analysis point of view, responses to the physical event are the actual signal: people report an event occurring in their vicinity, just like a sensor would (sample tweet: 'I think we just had an earthquake in Boston. Not kidding'). These responses are binary in nature and are comparable to threshold-based event detection in sensor networks (Abadi et al. 2005). In contrast to physical event reports, Twitter traffic addressing the cyber event

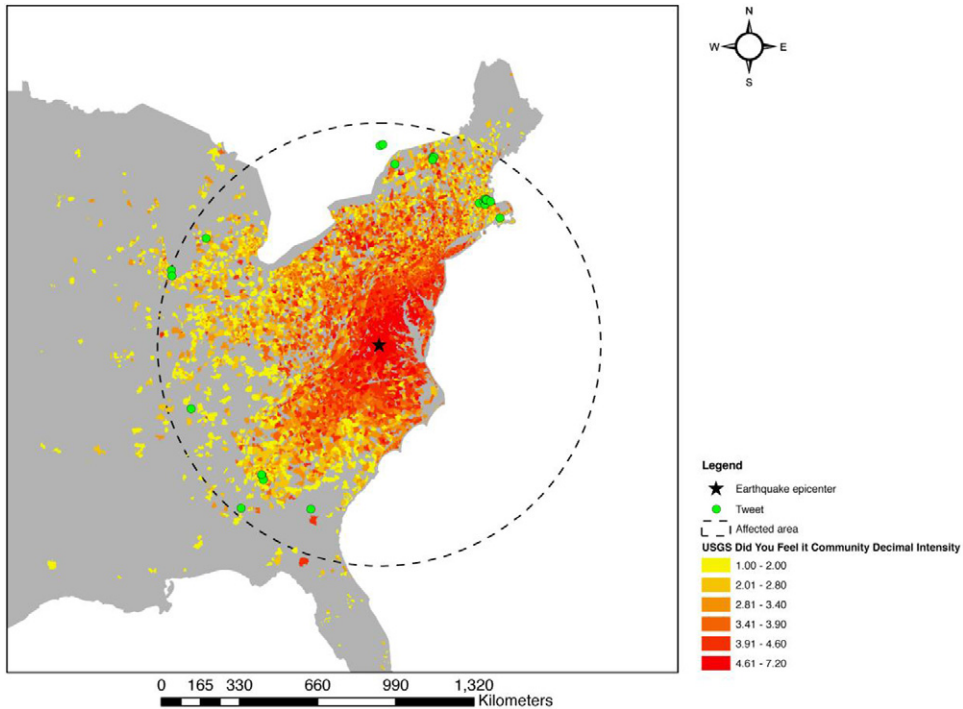


Figure 5 Locations of the 40 tweets in the shaded area of Figure 4 overlaid over the USGS CDI scale map. Tweet locations are marked as green circles. Color-coding in the graph is ranging from red (high perceived intensity) to yellow (lower perceived intensity). The dashed line shows a distance of approximately 950 km (8.5 degrees of angular distance) from the epicenter

(from outside the actual impact area) can be viewed as *noise* in the context of our analysis, as they are not reports of the event itself but reactions to the coverage of the event in the media (sample tweet: ‘Yep earthquake in dc/md area confirmed (source: family and friends VIA fb’). Such responses have a negative effect on attempts to localize the event using Twitter feeds. We can define a signal-to-noise ratio (SNR_T) metric over a time period T as:

$$SNR_T = \frac{\text{tweets within impact over time } T}{\text{total number tweets over time } T} \quad (1)$$

Therefore, a high SNR_T value would ensure better event localization potential. For our particular dataset the SNR_T value over the first 400 seconds was 95.8%. This is also visualized very effectively in the histogram of tweets shown in the bottom of Figure 4, where the values drop significantly after an angular distance of 8 degrees.

We can also observe in Figure 4 (top) a notable trend in response time within the impact area. Response time is represented by the lowest dots in each column of this figure, as they reflect the fastest response from a particular distance. We observe faster

responses closer to the epicenter, while the time of reaction drops almost linearly as we move away from it until we reach the edge of the impact area. This trend in reaction time within the impact area is consistent with the fact that the perceived local impact of an earthquake decreases as a function of distance, with tremors weakening away from the epicenter (Dengler and Dewey 1998). Beyond the impact area this trend ends, and even reverses, as is shown in the same figure, with reaction becoming faster as we move further towards the West Coast.

The earthquake is not an instantaneous event affecting all locations within its impact zone at the same time: the seismic waves propagate away from the epicenter, and the USGS provides estimates of the theoretical arrival times of the earthquake at different locations as a function of angular distance from the epicenter (http://neic.usgs.gov/neis/eq_depot/2011/eq_110823_se082311a/se082311a_t.html). Angular distance is calculated as the spherical central angle:

$$\Delta\hat{\sigma} = 2 \arcsin \left(\sqrt{\sin^2 \left(\frac{\phi_f - \phi_s}{2} \right) + \cos(\phi_s) \cos(\phi_f) \sin^2 \left(\frac{\Delta\lambda}{2} \right)} \right) \quad (2)$$

where ϕ_s and ϕ_f refer to the latitudes of the tweets and the epicenter, while $\Delta\lambda$ refers to the difference of their longitudes. For example, it is estimated that the wave reached Washington, D.C., within 24 seconds (at 17:51:28 UTC), Philadelphia within 49 seconds (at 17:51:53 UTC), New York City within 65 seconds (at 17:52:09 UTC), and further away it would have reached Los Angeles within 6 minutes 34 seconds (even though it would have been so weakened by then that it would have been undetectable by humans). Using this information, and in order to better understand the reaction time to the event we adjusted the data presented in Figure 4 to reflect response lag time compared to the local arrival time of the earthquake wave rather than the instance of the event at the epicenter. The adjusted results are shown in Figure 6. The axes are similarly defined to the ones we used in Figure 4 (top), but the data here are adjusted to correct for the arrival time of the seismic wave at each location. In Figure 6 we have also visualized by a light gray sliver the threshold zone between 700 and 950 km (approximately 6.5 and 8.5 angular degrees) that we used to generate the plot of Figure 5. We still see the very pronounced nearly linear increase in response time as we move away from the epicenter in Virginia towards the edges of the impact area. This is consistent with our earlier observation that the data may indicate that reaction time tends to be related to the perceived magnitude of the reported event.

The temporal reaction beyond the impact area is very interesting to observe in the adjusted data of Figure 6. We can see that the reaction time *outside* the impact area becomes negative as we consider areas like Salt Lake City, Phoenix, Los Angeles and Seattle. In these areas the Twitter community has already become aware of the event through the media *before* the actual event would have reached them. This observation shows the strong potential of using Twitter as an advanced warning system, as we see that it offers a lead time of one (e.g. Salt Lake City) to five (e.g. Los Angeles) minutes to alert the public at distant locations of a forthcoming incident.

If we consider responses over a longer period of time, say 60 minutes as shown in Figure 7, we can observe that similar to Figure 6 we still have the vast majority of tweets originating within the impact area. However, the SNR_T metric over this longer

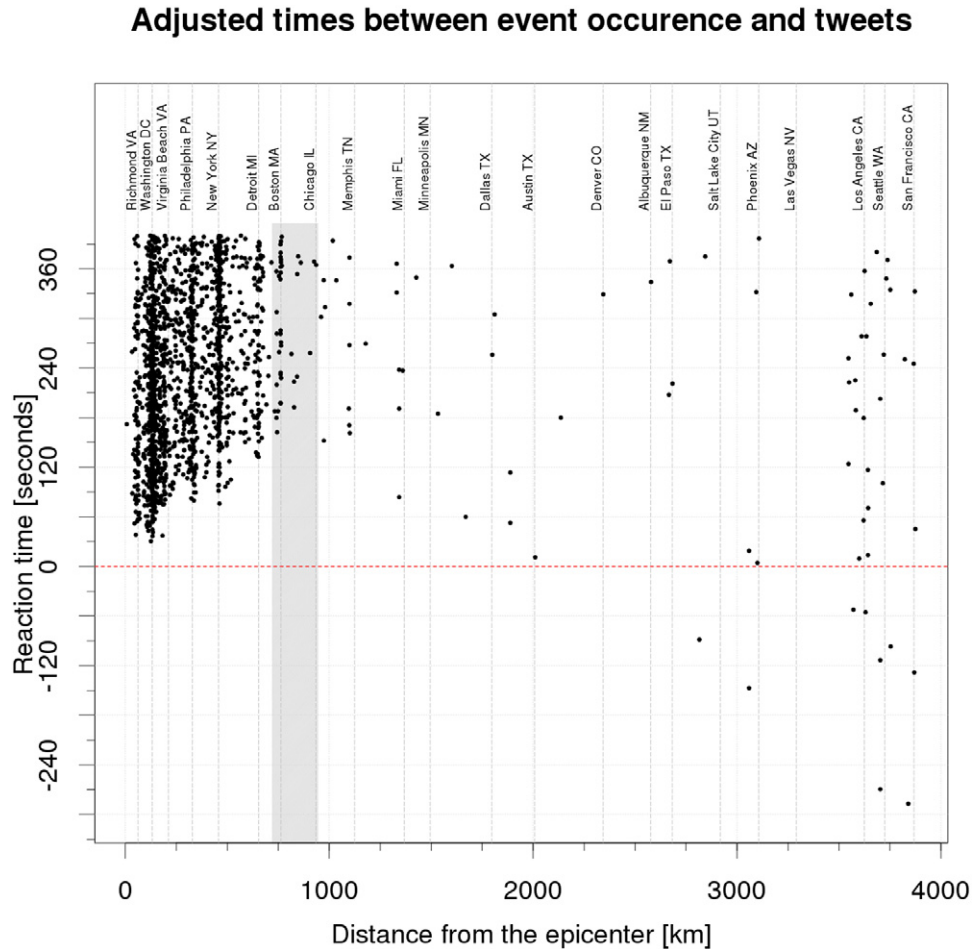


Figure 6 Response pattern as function of distance from epicenter: first 400 seconds. Reaction time is adjusted to reflect the local arrival time of the earthquake wave

period worsens, dropping to 88.1%. Nevertheless, we can easily identify the impact area in this figure too, as the left most dense portion of the graph extending towards the region of reversal in reaction times (Chicago). The data presented in Figure 7 are directly comparable to Figure 6, as they reflect the adjusted response time, to take into account the travel time of the seismic wave.

5.3 Event Coverage and Localization

In order to better visualize the geographic distribution of the Twitter-harvested earthquake reports we have mapped their locations at different time instances. In Figure 8 we show the evolution of the Twitter responses over the first five minutes of the event. The red dots indicate the locations from which we have received tweets reporting an earthquake for the corresponding time interval while the blue lines represent the seismic

Adjusted times between event occurrence and tweets




Figure 7 Response pattern as function of distance from epicenter for the first 60 minutes after the earthquake

wave at the start (nearest to Mineral, VA) and end of each analysis period. As we observe in Figure 8, already in the second analysis period (2 to 3 minutes) tweets appear ahead of the seismic wave. This demonstrates that as early as a couple of minutes after the event, Twittersphere awareness of the event overtakes the event's propagation in physical space.

We can see that reporting starts in the vicinity of the epicenter over the first two minutes, and starts spreading and densifying over the impact area as time passes by. We can see that the dots in the Twitter coverage maps as shown in Figure 8 correspond well to the high intensity regions of the CDI map of Figure 5. At the same time some sporadic responses start appearing, but we can see in the map that the major cluster of responses delineates very successfully the CDI impact area (compare Figures 5 and 8 lower right). This effect still holds as we monitor the period 5–10 minutes after the earthquake as seen in the left of Figure 9. By eight minutes after the event the seismic wave would have traveled across the U.S. as highlighted by the blue lines in the figure.

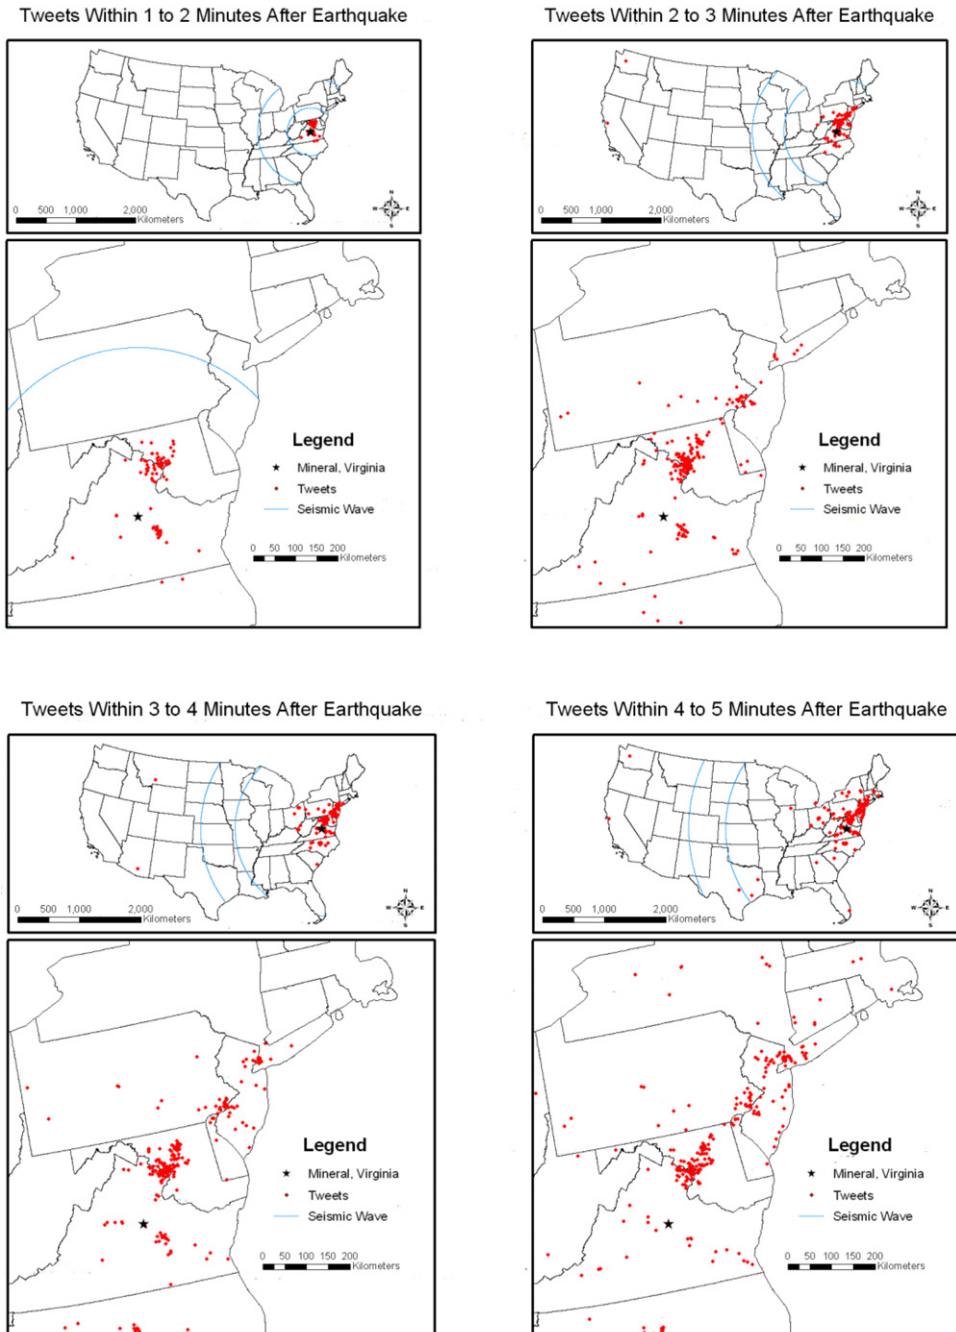


Figure 8 Mapping the originating locations of earthquake-reporting tweets at distinct time intervals during the first five minutes after the earthquake. The blue lines show earthquake wave propagation during the corresponding period

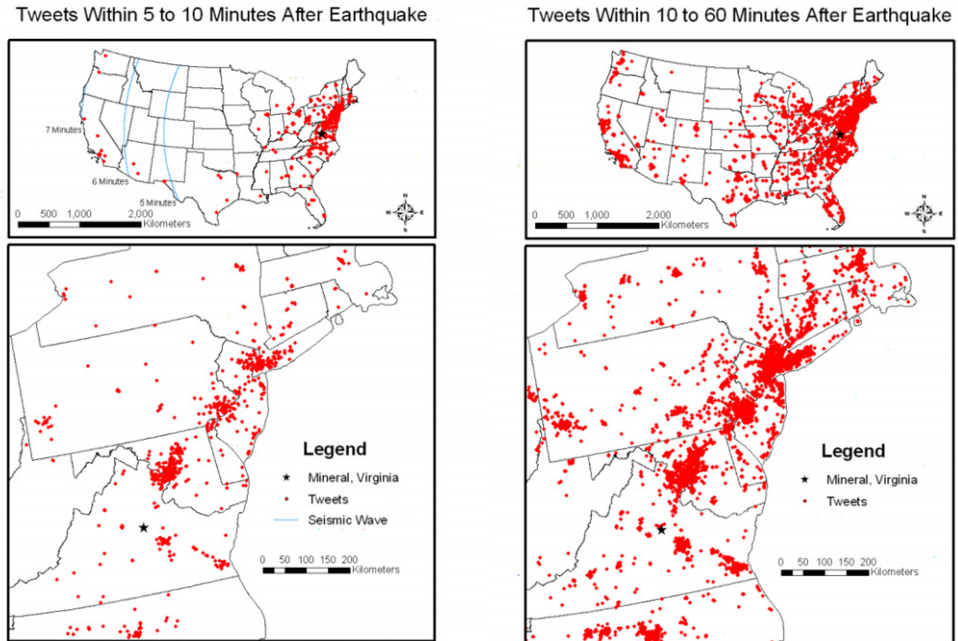


Figure 9 Mapping the originating locations of earthquake-reporting tweets 5–10 minutes (left) and 10–60 minutes (right) after the earthquake. The blue lines show earthquake wave propagation during the corresponding period

As we move beyond that time period, towards 60 minutes after the event as shown in Figure 9 (right) we see that the SNR_T values drop and response clusters also start forming outside the impact area. This is the effect of the increase in responses to the cyber event relative to responses to the physical event as we discussed in Section 5.2 above. Nevertheless, the East Coast impact area is still recognizable in the right of Figure 9. Based on these responses we can see that by focusing on Twitter feeds over the first few minutes after the event (up to 5 or even 10 minutes in this case) we can obtain a good approximation of the impact area, as responses to the cyber event remain relatively sparse during that early period.

5.4 Comparison of AGI to VGI Data Contributions

As we argued earlier, harvesting information from Twitter feeds represents a form of Ambient Geographic Information (AGI), which we consider to be a variant of what has traditionally been viewed as Volunteered Geographic Information (VGI) in terms of the purpose and process followed to contribute and collect this information. In order to assess the performance of VGI versus AGI we compare the data we harvested from Twitter to ‘Did you Feel It?’ the official VGI system of the USGS. Earlier, in Figures 5 and 8 we showed how well the two match in terms of geographic coverage. Here we present some additional analysis of the data contribution patterns in AGI and VGI.

In Figure 10 we show a comparison of the temporal rate of DYFI contributions versus geolocated earthquake references harvested from Twitter over a period of eight

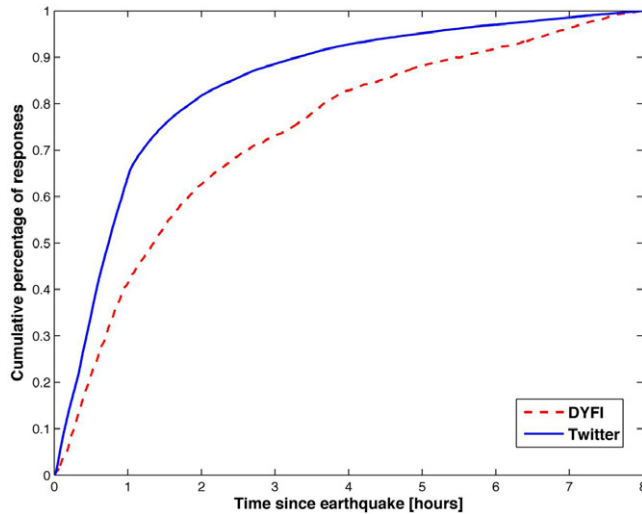


Figure 10 A comparison of the cumulative percentage of responses in DYFI and harvested geolocated tweets over the first eight hours since the earthquake

hours. We observe that the two curves are comparable, but Twitter shows a faster gain rate early on, and starts tapering off faster than DYFI, which displays a slower progression. In order to better compare the two curves we are displaying cumulative percentages of responses rather than actual numbers. The actual number of contributions to DYFI over the eight hour period is approximately 125,000 whereas our data corpus from a random 1% sample of Twitter includes approximately 23,500 tweets, and thus we could expect that the full Twitter traffic would correspond to nearly 2.3 million geolocated earthquake tweets over eight hours. This indicates that AGI flows faster and at higher volumes than VGI in this case.

In order to demonstrate the spatial variability of these contributions we visualize in Figures 11 and 12 patterns of DYFI contributions versus patterns of harvested geolocated tweets for two sample regions: Washington, D.C., and New York City. The data are aggregated over five-digit ZIP code regions, and we see that the patterns are indeed comparable. One issue of note is that the intensity scale between these two plots differs, as DYFI tends to consider as high values contributions at lower levels (e.g. 256–478) than the corresponding top tier of tweet streams (325–1,390). Nevertheless, top and low tiers in these two sets of maps correspond well, which indicates that the spatial variability of AGI and VGI contributions is comparable.

6 Discussion and Outlook

Social media feeds are rapidly emerging as a novel avenue for the dissemination of geographic information, as their content often includes references to events occurring at, or affecting specific locations. This Ambient Geographic Information (AGI) can be harvested from social media feeds and can be used to complement well-established data sources in providing us with advanced situational awareness for major physical events.

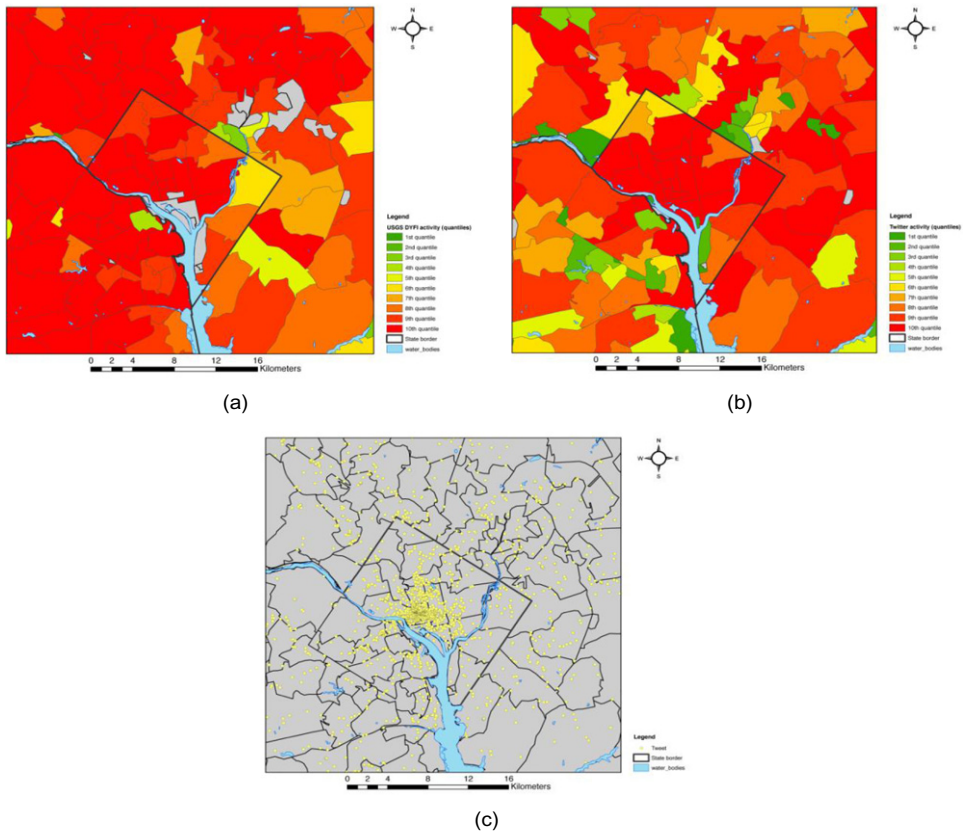


Figure 11 Washington, D.C.: (a) Plot of DYFI contributions aggregated per five digit ZIP code; (b) Plot of geolocated Twitter contributions aggregated per five digit ZIP code; (c) Plot of the geolocated Tweets

The objective of this article was to assess the quality of the content of such crowd-harvested information by considering the reaction of the Twitter community to a physical event with broad impact area, the Mineral, VA earthquake of August 23, 2011. We were particularly interested in assessing the role of these feeds as a novel type of sensor system for the detection and analysis of geographic events, and the experiments we present in this article support certain interesting observations.

The volume of geolocated information that can be harvested from Twitter feeds is very substantial. In our test case, using only a 1% sample of tweets we were able to collect approximately 100 accurately geolocated tweets within two minutes of the event, and nearly 1,000 such tweets within five minutes. Extrapolating from our 1% random sample to the complete Twitter traffic it is reasonable to expect close to 10,000 precisely geolocated tweet reports of the earthquake within two minutes of the event, and close to 100,000 within five minutes. This is a very substantial volume of reports, especially if we consider that the USGS DYFI website collected a comparable number of reports (125,000) over a period of eight hours. Based on these numbers we can argue that we harvest from Twitter and AGI a volume of reports comparable to the official

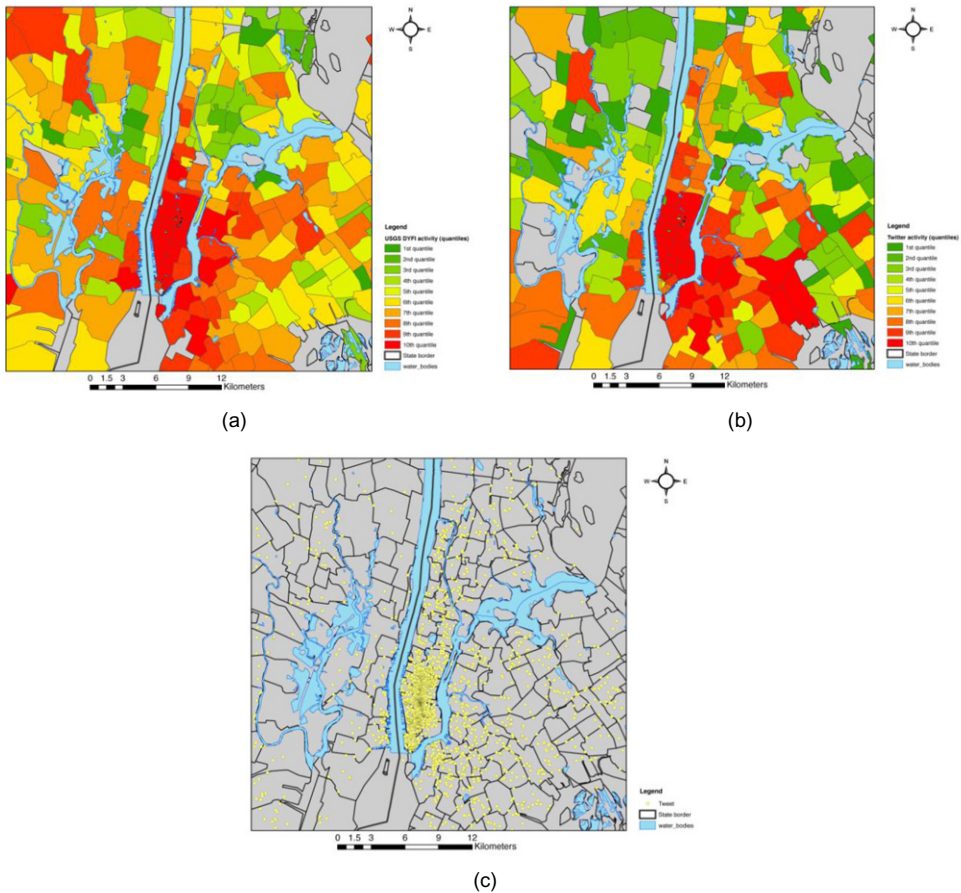


Figure 12 New York City: (a) Plot of DYFI contributions aggregated per five digit ZIP code; (b) Plot of geolocated Twitter contributions aggregated per five digit ZIP code; (c) Plot of the geolocated Tweets

VGI portal within 5% of the time. This reflects a substantial time gain, especially as it relates to disasters and the need to have timely information. In terms of sheer numbers, 100,000 reports is also considered a very substantial number if we consider that the planned USGS Advanced National Seismic System (ANSS) will comprise 7,000 earthquake monitoring units distributed all over the U.S. or that the EarthScope program which deploys a transportable array of 400 broadband seismometers over the U.S.

Moreover, the potential role of social media and AGI should not be underestimated as official sensor networks may be sparse or costly to implement, but human sensor networks are exponentially densifying with respect to the growing use of smart phones. They are not just dense but they are dense at the right places, namely where human populations are gathered, and arguably these are the places that are of the most interest under specific events. One might assume citizen reporting through Twitter will be less accurate than that of trained observers or actual recording stations and could be susceptible to false reports. However, authors such as Cox and Plale (2011) or Demirbas

et al. (2010) argue that such information can be accurate enough to improve localized weather reports. Which in a sense relates to the notion in the open source community that “given enough eyeballs, all bugs are shallow” (Raymond 1999) or the law of large numbers.

Another issue that needs to be considered is that while the data stream volume is impressive, we should not ignore the fact that the information conveyed by Twitter is qualitative but not quantitative. Unlike a seismometer, a tweet does not provide a metric of the magnitude of the earthquake but rather conveys binary information: that an earthquake was felt at a given location. However, the massive amount of data harvested from Twitter allows us to overcome this limitation, and identifies the extent of the impact area of the earthquake event as shown in Figures 8 and 9.

We also observed interesting spatial and temporal trends in Twitter reports of the earthquake event. More specifically, the majority of tweets originated from within the impact area, and slowly over time diffused across the country as we move from tweeting about the event to tweeting about news of the event. Accordingly, an early collection of tweets (e.g. in the first 10 minutes after the event) can be used to provide a fast and good approximation of the impact area (especially if the context of the tweet and associated images were to be analyzed as well), which is valuable information for response and recovery operations. Furthermore, reaction time tends to increase away from the epicenter, which may serve as an indicator to estimate the event epicenter simply by considering the temporal stamps of relevant Twitter data streams. Moreover, one could also explore a “social epicenter” or “cyber epicenter”, in the sense that the central point of the cluster of messages about the phenomena could be delineated and this would likely move in time, as more people report the event. Our dataset also showed that information dissemination through Twitter can travel faster than the physical event to distant locations, as demonstrated vividly by the negative reaction times of the West Coast in Figure 5. This is a very critical observation supporting the use of social media in general, and Twitter in particular, for providing novel geographic information and thus acting as an early warning system for large-scale incidents. Together with the abovementioned ability to obtain an accurate assessment of the impact area, this supports the notion that by harvesting geospatial information from Twitter feeds we can rapidly gain valuable information on the impact of a physical event. Of course in the experiments we showed here the data were collected after the event, but nevertheless we have demonstrated through our work that tweet content enables the rapid identification of impact areas and enhances our situational awareness. In order to fully realize this potential we need to take advantage of systems that perform real-time trend detection over the Twitter stream (e.g. Mathioudakis and Koudas 2010). Recently, Zhao et al. (2011) reported on the use of an adaptive sliding window application of a lexicon-based content analysis solution that is able to detect events in Twitter within 40 seconds of the actual event. The integration of such systems with the type of analysis we demonstrated here supports the emergence of Twitter as an effective distributed sensor system for event detection and impact assessment.

The geographic focus of our work complements earlier work on data mining approaches that addressed disease outbreak detection by harvesting information from news articles (Brownstein et al. 2009), Internet search engines (e.g. Polgreen et al. 2008), blogs at large (Corley et al. 2010), and Twitter specifically (Ritterman et al. 2009, Culotta 2010). The observations we made in this article highlight the very rich potential offered by harvesting ambient geographic information from Twitter to

enhance our situational awareness capabilities, and the evolution of geographic information from traditional formats (e.g. maps and imagery) and authoritative sources to distributed content provided by the general public.

Acknowledgments

We would like to acknowledge Eric Fischer and Maksim Tsvetovat for their assistance.

Note

- 1 Twitter provides free access to a 1% sample of all tweets via their API and permits the geolocation analysis of such harvested information. For this analysis, the 1% sample was based on keywords and hashtags relating to earthquakes.

References

- Abadi D, Madden S, and Lindner W 2005 REED: Robust, Efficient Filtering and Event Detection in sensor networks. In *Proceedings of the Thirty-first International Conference on Very Large Data Bases (VLDB'05)*, Trondheim, Norway: 769–80
- Akyildiz I F, Sankarasubramanian Y, and Cayirci E 2002 Wireless sensor networks: A survey. *Computer Networks* 38: 393–422
- Atkinson G M and Wald D 2007 “Did You Feel It?” intensity data: A surprisingly good measure of earthquake ground motion. *Seismological Research Letters*, 78: 362–68
- Backstrom L, Sun E, and Marlow C 2010 Find me if you can: Improving geographical prediction with social and spatial proximity. In *Proceedings of the Nineteenth International World Wide Web Conference (WWW'10)*, Raleigh, North Carolina: 61–70
- Brownstein J S, Clark C, Freitfeld B S, and Madoff L C 2009 Digital disease detection: Harnessing the Web for public health surveillance. *New England Journal of Medicine* 360: 2153–57
- Burke J, Estrin D, Hansen M, Parker A, Ramanathan N, Reddy S, and Srivastava M B 2006 Participatory sensing. In *Proceedings of the Workshop on the World-Sensor-Web (WSW'06), Mobile Device Centric Sensor Networks and Applications*, Boulder, Colorado: 117–34
- Cheng Z, Caverlee J, and Lee K 2010 You are where you Tweet: A content-based approach to geo-locating Twitter users. In *Proceedings of the ACM Conference on Information and Knowledge Management (CIKM'10)*, Toronto, Ontario: 759–68
- Christakis N A and Fowler J H 2010 Social network sensors for early detection of contagious outbreaks. *PLoS ONE* 5 (9): e12948
- Christensen C 2011 Twitter revolutions? Addressing social media and dissent. *The Communication Review* 14: 155–57
- Christian Science Monitor 2011 Social Media Drive Occupy Wall Street: Do They Also Divulge Its Secrets? WWW document, <http://www.csmonitor.com/USA/Politics/2011/1012/Social-media-drive-Occupy-Wall-Street.-Do-they-also-divulge-its-secrets>
- Corley C D, Cook D J, Mikler A R, and Singh K P 2010 Text and structural data mining of influenza mentions in web and social media. *International Journal of Environmental Research and Public Health* 7: 596–615
- Cox J and Plale B 2011 *Improving Automatic Weather Observations with the Public Twitter Stream*. Bloomington, IN, Indiana University, School of Informatics and Computing Research Report (available at: <http://www.cs.indiana.edu/Research/techreports/>)

- Culotta A 2010 'Towards detecting Influenza Epidemics by Analyzing Twitter Messages', *Proceedings of the First Workshop on Social Media Analytics*, Washington, D.C., pp. 115–122
- De Longueville B, Smith R S, and Luraschi G 2009 "OMG, from here, I can see the flames!": A use case of mining location based social networks to acquire spatio-temporal data on forest fires. In *Proceedings of the International Workshop on Location-Based Social Networks (LBSN '09)*, Seattle, Washington: 73–80
- Demirbas M, Bayir M A, Akcora C G, Yilmaz Y S, and Ferhatosmanoglu H 2010 Crowd-sourced sensing and collaboration using Twitter. In *Proceedings of the IEEE International Symposium on a World of Wireless, Mobile and Multimedia Networks (WOWMOM)*, Montreal, Quebec: 1–9
- Dengler L A and Dewey J W 1998 An intensity survey of households affected by the Northridge, California, earthquake of 17 January 1994. *Bulletin of the Seismological Society of America* 88: 441–62
- Doty N and Wilde E 2010 Geolocation privacy and application platforms. In *Proceedings of the Third ACM SIGSPATIAL International Workshop on Security and Privacy in GIS and LBS*, San Jose, California: 65–69
- Dyregrov A, Solomon R, and Bassoe C F 2000 Mental mobilization process in critical incident stress situations. *International Journal of Emergency Mental Health* 2(2): 73–81
- Earle P, Guy M, Buckmaster R, Ostrum C, Horvath S, and Vaughan A 2010 OMG earthquake! Can Twitter improve earthquake response? *Seismological Research Letters* 81: 246–51
- Eriksson B, Barford P, Sommers J, and Nowak R 2010 A learning-based approach for IP geolocation. In Krishnamurthy A and Plattner B (eds) *Passive and Active Measurement*. Berlin, Springer Lecture Notes in Computer Science Vol. 6032: 171–80
- Field K and O'Brien J 2010 Cartoblography: Experiments in using and organizing the spatial context of micro-blogging. *Transactions in GIS* 15: 5–23
- Fritz S, MacCallum I, Schill C, Perger C, Grillmayer R, Achard F, Kraxner F, and Obersteiner M 2009 Geo-Wiki.Org: The use of crowdsourcing to improve global land cover. *Remote Sensing*, 1: 345–54
- Ginsberg J, Mohebbi M H, Patel R S, Brammer L, Smolinski M S, and Brilliant L 2009 Detecting influenza epidemics using search engine query data. *Nature* 457: 1012–14
- Goodchild M F 2007 Citizens as sensors: The world of volunteered geography. *GeoJournal* 69: 211–21
- Goodchild M F and Glennon J A 2010 Crowdsourcing geographic information for disaster response: A research frontier. *International Journal of Digital Earth* 3: 231–41
- Hecht B, Hong L, Suh B, and Chi E 2011 Tweets from Justin Bieber's Heart: The dynamics of the 'location' field in user profiles. In *Proceedings of the ACM Conference on Human Factors in Computing Systems*, Vancouver, British Columbia: 237–46
- Honey C and Herring S C 2009 Beyond microblogging: Conversation and collaboration via Twitter. In *Proceedings of the Forty-Second Hawaii International Conference on System Sciences (HICSS-42)*, Big Island, Hawaii: 1–10
- Howard P N, Agarwal S D, and Mazammil M 2011 When do states disconnect their digital networks? Responses to the political uses of social media. *The Communication Review* 14: 216–32
- Howe J 2008 *Crowdsourcing: Why the Power of the Crowd is Driving the Future of Business*. New York, Crown Business
- Hyvarinen O and Saltikoff E 2010 Social media as a source of meteorological observations. *Monthly Weather Review* 138: 3175–84
- Huffington Post 2011 Occupy Wall Street: Social Media's Role in Social Change. WWW document, http://www.huffingtonpost.com/2011/10/06/occupy-wall-street-socialmedia_n_999178.html
- Java A, Song X, Finin T, and Tseng B 2009 Why we Twitter: An analysis of a microblogging community. In Goebel R, Siekmann J, and Wahlster W (eds) *Advances in Web Mining and Web Usage Analysis*. Berlin, Springer Lecture Notes in Computer Science Vol. 5439: 118–38
- Kaplan A M and Haenlein M 2010 Users of the World unite! The challenges and opportunities of social media. *Business Horizons* 53: 59–68
- Kosala R and Blockeel H 2000 Web mining research: A survey. *SIGKDD Explorations* 2: 1–15

- Ludford P J, Priedhorsky R, Reily K, and Terveen L G 2007 Capturing, sharing, and using local place information. In *Proceedings of the ACM Conference on Human Factors in Computing Systems (CHI '07)*, San Jose, California: 1235–44
- MacEachren A M, Jaiswal A, Robinson A C, Pezanowski S, Savelyev A, Mitra P, Zhang X, and Blanford J 2011 SensePlace2: GeoTwitter analytics support for situational awareness. In *Proceedings of the IEEE Conference on Visual Analytics Science and Technology*, Providence, Rhode Island
- Mathioudakis M and Koudas N 2010 TwitterMonitor: Trend detection over the Twitter stream. In *Proceedings of the International ACM Conference on Management of Data (SIGMOD'10)*, Indianapolis, Indiana: 1155–58
- Miller C and Wortham J 2010 Technology Aside, Most People Still Decline to Be Located. *New York Times*. (30 August; available at: <http://www.nytimes.com/2010/08/30/technology/30location.html?src=tp>)
- Naaman M, Boase J, and Lai C-H 2010 'Is it really about me? Message content in social awareness streams'. In *Proceedings of the 2010 ACM Conference on Computer Supported Cooperative Work – (CSCW'10)*, Savannah, Georgia: 189–92
- Nardi B A, Schiano D J, Gumbrecht M, and Schwartz L 2004 Why we blog. *Communications of the ACM* 47(12): 41–6
- Nessler D, Mecklinger A, and Penney T B 2005 Perceptual fluency, semantic familiarity and recognition-related familiarity: An electrophysiological exploration. *Cognitive Brain Research* 22: 265–88
- Newsweek 2009 A Twitter Timeline of the Iran Election. *Newsweek* (available at <http://www.newsweek.com/2009/06/25/a-twitter-timeline-of-the-iran-election.html>)
- Nittel S, Labrinidis A, and Stefanidis A (eds) 2008 *Advances in Geosensor Networks*. Berlin, Springer Lecture Notes in Computer Science Vol. 4540
- Okasaki M and Matsuo Y 2012 Semantic Twitter: Analyzing tweets for real-time event notification. In Breslin J G, Burg T N, Kim H-G, Raftery T, and Schmidt J-H (eds) *Recent Trends and Development in Social Software*. Berlin, Springer Lecture Notes in Computer Science Vol. 6045: 63–74
- Poese I, Uhlig S, Kaafar M A, Donnet B, and Gueye B 2011 IP Geolocation databases: Unreliable? *Computer Communication Review* 4 (2): 53–56
- Polgreen P M, Chen Y, Pennock D M and Nelson F D 2008 Using Internet searches for influenza surveillance. *Clinical Infectious Diseases*, 47: 1443–48
- Popescu A and Grefenstette G 2010 Mining user home location and gender from Flickr tags. In *Proceedings of the Fourth International Conference on Weblogs and Social Media – (ICWSM '10)*, Washington, D.C.: 307–10
- Raymond E S 1999 *The Cathedral and the Bazaar: Musings on Linux and Open Source by an Accidental Revolutionary*. North Sebastopol, CA, O'Reilly Media
- Russell M A 2011 *Mining the Social Web*. North Sebastopol, CA, O'Reilly Media
- Ritterman J, Osborne M, and Klein E 2009 Using prediction markets and Twitter to predict a swine flu pandemic. In *Proceedings of the First International Workshop on Mining Social Media*, Sevilla, Spain: 9–18
- Sakaki T, Okazaki M, and Matsuo Y 2010 Earthquake shakes Twitter users: Real-time event detection by social sensors. In *Proceedings of the International Conference on the World Wide Web (WWW '10)*, Raleigh, North Carolina: 851–60
- Selavo L, Wood A, Cao Q, Sookoor T, Liu H, Srinivasan A, Wu Y, Kang W, Stankovic J, Young D, and Porter J 2007 LUSTER: Wireless sensor network for environmental research. In *Proceedings of the Fifth International Conference on Embedded Networked Sensor Systems*, Sydney, Australia: 103–16
- Singh V K, Gao M, and Jain R 2010 Situation detection and control using spatio-temporal analysis of microblogs. In *Proceedings of the International Conference on the World Wide Web (WWW '10)*, Raleigh, North Carolina: 1181–82
- Stefanidis A, Crooks A, and Radzikowski J 2012 Harvesting ambient geospatial information from social media feeds. *GeoJournal* 77: in press
- Steinfeld A, Zimmerman J, Tomasic A, Yoo A, and Aziz R 2011 Mobile transit rider information via universal design and crowdsourcing. In *Proceedings of the 2011 Annual Meeting of the Transportation Research Board*, Washington, D.C.

- Sternberg R J 2009 *Cognitive Psychology*. Florence, KY, Wadsworth Publishing
- Vieweg S, Hughes A L, Starbird K, and Palen L 2010 Microblogging during two natural hazards events: What Twitter may contribute to situational awareness. In *Proceedings of the ACM Conference on Computer Human Interaction (CHI '10)*, Atlanta, Georgia: 1079–88
- Wald D J, Quitoriano V, Dengler L A, and Dewey J W 1999 Utilization of the Internet for rapid community intensity maps. *Seismological Research Letters* 70: 680–93
- Yamaguchi S, Hale L A, D'Esposito M, and Knight R T 2004 Rapid prefrontal-hippocampal habituation to novel events. *Journal of Neuroscience* 24: 5356–63
- Zhao D and Rosson M B 2009 How and why people Twitter: The role that micro-blogging plays in informal communication at work. In *Proceedings of the ACM International Conference on Supporting Group Work*, Sanibel Island, Florida: 243–52
- Zhao S, Zhong L, Wickramasuriya J, and Vasudevan V 2011 Human as real-time sensors of social and physical events: A case study of Twitter and sports games. Houston, TX, Rice University Technical Report No. TR0620